

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, recall_score, f1_score, confusion_matrix
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/diabetes2.csv')
df.sample(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>631</b>	0	102	78	40	90	34.5	0.238	24	0
<b>411</b>	1	112	72	30	176	34.4	0.528	25	0
<b>640</b>	0	102	86	17	105	29.3	0.695	27	0
<b>16</b>	0	118	84	47	230	45.8	0.551	31	1
<b>638</b>	7	97	76	32	91	40.9	0.871	32	1

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
df.isnull().sum()
```

	0
<b>Pregnancies</b>	0
<b>Glucose</b>	0
<b>BloodPressure</b>	0
<b>SkinThickness</b>	0
<b>Insulin</b>	0
<b>BMI</b>	0
<b>DiabetesPedigreeFunction</b>	0
<b>Age</b>	0
<b>Outcome</b>	0

**dtype:** int64

df.duplicated().sum()

np.int64(0)

```
import seaborn as sns

sns.set(style="whitegrid")

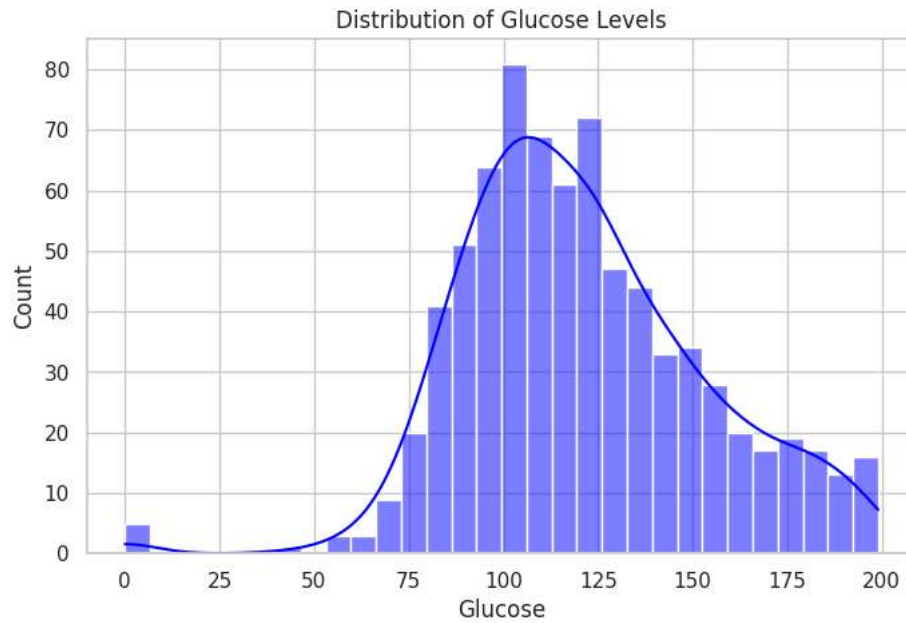
# Distribution of Glucose levels
plt.figure(figsize=(8,5))
sns.histplot(df['Glucose'], kde=True, bins=30, color="blue")
plt.title("Distribution of Glucose Levels")
plt.xlabel("Glucose")
plt.ylabel("Count")
plt.show()

# Outcome count (0 = Non-diabetic, 1 = Diabetic)
plt.figure(figsize=(6,4))
sns.countplot(x='Outcome', data=df, palette='Set2')
plt.title("Diabetes Outcome Count")
plt.xlabel("Outcome (0 = No, 1 = Yes)")
plt.ylabel("Count")
plt.show()

# Correlation heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()

# Age distribution by outcome
plt.figure(figsize=(8,5))
sns.boxplot(x="Outcome", y="Age", data=df, palette="Set3")
plt.title("Age Distribution by Diabetes Outcome")
plt.show()
```







/tmp/ipython-input-1448572228.py:15: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Outcome', data=df, palette='Set2')
```



```
X_train, X_test, y_train, y_test = train_test_split(df.drop(['Outcome'], axis=1),
                                                    df['Outcome'],
                                                    test_size=0.2,
                                                    random_state=2)
```

```
std = StandardScaler()
```

```
X_train = std.fit_transform(X_train)
X_test = std.transform(X_test)
```

```
model = LogisticRegression(max_iter=1000, class_weight='balanced')
```

```
model.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression(class_weight='balanced', max_iter=1000)
```

```
y_pred = model.predict(X_test)
```

```
print('Accuracy: ', accuracy_score(y_test, y_pred))
print('Recall: ', recall_score(y_test, y_pred))
print('f1 score: ', f1_score(y_test, y_pred))
print('Confusion Matrix: ', confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.7532467532467533
```

```
f1 score: 0.6346153846153846
```

```
Confusion Matrix: [[12 33]]
```

```
FutureWarning: Assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.
```

```
sns.boxplot(x="Outcome", y="Age", data=df, palette="Set3")
```

```

import numpy as np
import matplotlib.pyplot as plt

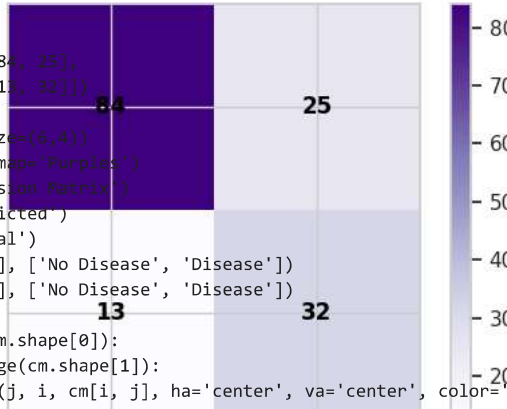
cm = np.array([[84, 25],
               [13, 32]])

plt.figure(figsize=(6,4))
plt.imshow(cm, cmap='Purples')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks([0, 1], ['No Disease', 'Disease'])
plt.yticks([0, 1], ['No Disease', 'Disease'])

for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        plt.text(j, i, cm[i, j], ha='center', va='center', color='black', fontsize=12, fontweight='bold')

plt.colorbar()
plt.show()

```



The confusion matrix shows the following counts:

	Actual No Disease	Actual Disease
Predicted No Disease	84	25
Predicted Disease	13	32

```

from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))

```

	precision	recall	f1-score	support
0	0.87	0.76	0.81	109
1	0.56	0.73	0.63	45
accuracy			0.75	154
macro avg	0.72	0.75	0.72	154
weighted avg	0.78	0.75	0.76	154

Start coding or [generate](#) with AI.