





```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, recall_score, f1_score, confusion_matrix
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('/content/diabetes2.csv')
df.sample(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
655	2	155	52	27	540	38.7	0.240	25	1	
253	0	86	68	32	0	35.8	0.238	25	0	
404	5	168	64	0	0	32.9	0.135	41	1	
479	4	132	86	31	0	28.0	0.419	63	0	
710	3	158	64	13	387	31.2	0.295	24	0	

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000	

```
df.isnull().sum()
```

	0
Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

dtype: int64

df.duplicated().sum()

np.int64(0)

```
print((df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] == 0).sum())
```

```
df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']] = df[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']].replace(0, pd.NaT)
df.fillna(df.mean(), inplace=True)
```

Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0

dtype: int64

```
X_train, X_test, y_train, y_test = train_test_split(df.drop(['Outcome'], axis=1),
                                                    df['Outcome'],
                                                    test_size=0.2,
                                                    random_state=2)
```

```
std = StandardScaler()
```

```
X_train = std.fit_transform(X_train)
X_test = std.transform(X_test)
```

```
model = LogisticRegression(max_iter=1000, class_weight='balanced')
```

```
model.fit(X_train, y_train)
```

```
LogisticRegression  
LogisticRegression(class_weight='balanced', max_iter=1000)
```

```
y_pred = model.predict(X_test)
```


```
print('Accuracy: ', accuracy_score(y_test, y_pred))  
print('Recall: ', recall_score(y_test, y_pred))  
print('f1 score: ', f1_score(y_test, y_pred))  
print('Confusion Matrix: ', confusion_matrix(y_test, y_pred))
```

```
Accuracy: 0.7532467532467533  
Recall: 0.7111111111111111  
f1 score: 0.6274509803921569  
Confusion Matrix: [[84 25]  
 [13 32]]
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(6,4))  
plt.imshow(cm, cmap='Purples')  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.xticks([0, 1], ['No Disease', 'Disease'])  
plt.yticks([0, 1], ['No Disease', 'Disease'])
```

```
([<matplotlib.axis.YTick at 0x7ff1876fe030>,  
 <matplotlib.axis.YTick at 0x7ff18c1fc890>],  
 [Text(0, 0, 'No Disease'), Text(0, 1, 'Disease')])
```



```
from sklearn.metrics import classification_report  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.77	0.82	109
1	0.56	0.71	0.63	45
accuracy			0.75	154
macro avg	0.71	0.74	0.72	154
weighted avg	0.78	0.75	0.76	154